


Take-Home Practice on Neural Networks (Non-Graded) -- Solutions

1. Backpropagation example with 2 inputs, 2 hidden units and 2 outputs



Assume that the neurons have sigmoid activation function, learning rate of 1 and answer the following questions:

- Perform a forward pass on the network and determine the outputs at E and F.
- Perform back-propagation on the output layer neurons($\text{target}_E=1$ and $\text{target}_F=0$) and determine the updated weights for w_5 , w_6 , w_7 and w_8 .
- Perform back-propagation on the hidden layer neurons and determine the updated weights for w_1 , w_2 , w_3 , and w_4 .

Useful sigmoid values:

x	0.175	0.4	0.42	0.5	0.6	0.73	0.82
sigmoid(x)	0.543	0.598	0.6	0.622	0.645	0.67	0.694

Solution 1a)

$$\text{in}_C = (0.25 \cdot 0.6) + (0.5 \cdot 0.9) = 0.15 + 0.45 = 0.6$$

$$\text{out}_C = \text{sigmoid}(0.6) = 0.645$$

$$\text{in}_D = (0.25 \cdot 0.1) + (0.5 \cdot 0.3) = 0.025 + 0.15 = 0.175$$

$$\text{out}_D = \text{sigmoid}(0.175) = 0.543$$

$$\text{in}_E = (0.645 \cdot 0.4) + (0.543 \cdot 0.3) = 0.258 + 0.162 = 0.42$$

$$\text{out}_E = \text{sigmoid}(0.42) = 0.6$$

$$\text{in}_F = (0.645 \cdot 0.2) + (0.543 \cdot 0.5) = 0.129 + 0.271 = 0.4$$

$$\text{out}_F = \text{sigmoid}(0.4) = 0.598$$

Solution 1b)

$$\delta_E = out_E(1 - out_E)(target_E - out_E) = 0.6(1 - 0.6)(1 - 0.6) = 0.096$$

$$\delta_F = out_F(1 - out_F)(target_F - out_F) = 0.598(0 - 0.598)(1 - 0.598) = -0.14$$

$$w_5^* = w_5 + \eta \cdot \delta_E \cdot out_C = 0.4 + 1 * (0.096 \cdot 0.645) = 0.46$$

$$w_6^* = w_6 + \eta \cdot \delta_F \cdot out_C = 0.2 + 1 * (-0.14 \cdot 0.645) = 0.11$$

$$w_7^* = w_7 + \eta \cdot \delta_E \cdot out_D = 0.3 + 1 * (0.096 \cdot 0.543) = 0.352$$

$$w_8^* = w_8 + \eta \cdot \delta_F \cdot out_D = 0.5 + 1 * (-0.14 \cdot 0.543) = 0.42$$

Solution 1c)

$$\delta_C = out_C(1 - out_C)(\delta_E \cdot w_5 + \delta_F \cdot w_6) = 0.645(1 - 0.645)(0.096 \cdot 0.46 + -0.14 \cdot 0.11) = 0.007$$

$$\delta_D = out_D(1 - out_D)(\delta_E \cdot w_7 + \delta_F \cdot w_8) = 0.543(1 - 0.543)(0.096 \cdot 0.352 + -0.14 \cdot 0.43) = -0.007$$

$$w_1^* = w_1 + \eta \cdot \delta_C \cdot in_A = 0.6 + 1 * (0.007 \cdot 0.25) = 0.601$$

$$w_2^* = w_2 + \eta \cdot \delta_D \cdot in_A = 0.1 + 1 * (-0.007 \cdot 0.25) = 0.09825$$

$$w_3^* = w_3 + \eta \cdot \delta_C \cdot in_B = 0.9 + 1 * (0.007 \cdot 0.5) = 0.904$$

$$w_4^* = w_4 + \eta \cdot \delta_D \cdot in_B = 0.3 + 1 * (-0.007 \cdot 0.5) = 0.297$$

2. The following matrix represents the weights of a hopfield network with the vectors $V_1(0, 1, 0, 1)$, $V_2(1, 0, 0, 1)$ stored. Use this matrix and answer the following questions.

$$W = \begin{bmatrix} 0 & -2 & 0 & 0 \\ -2 & 0 & 0 & 0 \\ 0 & 0 & 0 & -2 \\ 0 & 0 & -2 & 0 \end{bmatrix}$$

- What is the weight matrix when a new vector $V_3(0, 1, 1, 0)$ is added to this network?
- Use the weight matrix obtained in part a, and assume that the order of node updates is 2, 3, 4, 1. What memory does the network converge to if V_{in} is $(1, 0, 1, 0)$? (Show the input vector after each update and the final attractor that the network converges to)

Solution 2a):

The weight matrix for storing just V_3 is:

$$W_3 = \begin{bmatrix} 0 & -1 & -1 & 1 \\ -1 & 0 & 1 & -1 \\ -1 & 1 & 0 & -1 \\ 1 & -1 & -1 & 0 \end{bmatrix}$$

This is calculated using $W_3^{ij} = (2*V_3^i - 1) \times (2*V_3^j - 1)$.

The final weight matrix $W = W + W_3$

$$W = \begin{bmatrix} 0 & -3 & -1 & 1 \\ -3 & 0 & 1 & -1 \\ -1 & 1 & 0 & -3 \\ 1 & -1 & -3 & 0 \end{bmatrix}$$

Solution2b):**Iteration 1:**

Node updates order: [2, 3, 4, 1]

Updating node 2:

$$V_{2in}: W[2] * V_2 = (-3, 0, 1, -1) * (1, 0, 1, 0) = -2 < 0 \Rightarrow V_{2in} = 0 \text{ (didn't change)}, V_{in} = (1, 0, 1, 0)$$

Updating node 3:

$$V_{3in}: (-1, 1, 0, -3) * (1, 0, 1, 0) = -1 < 0 \Rightarrow V_{3in} = 0 \text{ (changed)}, V_{in} = (1, 0, 0, 0)$$

Updating node 4:

$$V_{4in}: (1, -1, -3, 0) * (1, 0, 0, 0) = 1 \geq 0 \Rightarrow V_{4in} = 1 \text{ (changed)}, V_{in} = (1, 0, 0, 1)$$

Updating node 1:

$$V_{1in}: (0, -3, -1, 1) * (1, 0, 0, 1) = 1 \geq 0 \Rightarrow V_{1in} = 1 \text{ (didn't change)}, V_{in} = (1, 0, 0, 1)$$

Since there were updates made to the input vector, we need to perform another iteration to test for convergence.

Iteration 2:

Node updates order: [2, 3, 4, 1]

Updating node 2:

$$V_{2in}: (-3, 0, 1, -1) * (1, 0, 0, 1) = -4 < 0 \Rightarrow V_{2in} = 0 \text{ (didn't change)}, V_{in} = (1, 0, 0, 1)$$

Updating node 3:

$V_{3in}: (-1, 1, 0, -3) * (1, 0, 0, 1) = -4 < 0 \Rightarrow V_{3in} = 0$ (didn't change), $V_{in} = (1, 0, 0, 1)$

Updating node 4:

$V_{4in}: (1, -1, -3, 0) * (1, 0, 0, 1) = 1 \geq 0 \Rightarrow V_{4in} = 1$ (didn't change), $V_{in} = (1, 0, 0, 1)$

Updating node 1:


$V_{1in}: (0, -3, -1, 1) * (1, 0, 0, 1) = 1 \geq 0 \Rightarrow V_{1in} = 1$ (didn't change), $V_{in} = (1, 0, 0, 1)$

Since there were no updates made to the input vector, we conclude that the network converged to the stored memory $V_2(1, 0, 0, 1)$.

3. Consider the input vectors

$I_1 = (1.1, 1.7, 1.8)$, $I_2 = (0, 0, 0)$, $I_3 = (0, 0.5, 1.5)$, $I_4 = (1, 0, 0)$, $I_5 = (0.5, 0.5, 0.5)$, $I_6 = (1, 1, 1)$.

We are using a 3 node Self-Organizing Map network with initial weights $W(0)$ as shown below:



Using the above information, answer the following questions:

- If the neighborhood radius is $R = 1$ (ie. we consider neighbors at distance 1 of a node when updating weights), learning rate $\eta = 0.5$, and we consider the inputs in order (I_1 to I_6), **what is the weight matrix after the first epoch** (after all inputs are presented once)?
- What clusters (A, B or C) are the inputs assigned to at the end of the first epoch?
- Assuming a geometric decrease in learning rate of 0.5, what is the learning rate for the second and third epochs?

Solution 3a):

For the first input $I_1(1.1, 1.7, 1.8)$:

$$\begin{aligned} \text{Euclidean Distance of A from } I_1 &= D^2(A, I_1) = \sum_{j=1 \text{ to } 3} (I_1^{(j)} - W_A^{(j)})^2 \\ &= (1.1 - 0.2)^2 + (1.7 - 0.7)^2 + (1.8 - 0.3)^2 = 4.1 \end{aligned}$$

Similarly, $D^2(B, I_1) = 4.4$ and $D^2(C, I_1) = 1.1$

Since C has the shortest distance to I_1 , we pick C as the winning neuron, and update the weights of C and all the neurons at a distance of 1 (since $R=1$) from it, which is only neuron A, thus

We update W_C and W_A :

$$W_{C1} = W_{C1} + \eta * (I_{11} - W_{C1})) = 1 + 0.5 * (1.1 - 1) = 1.05$$

$$W_{C2} = W_{C2} + \eta * (I_{12} - W_{C1})) = 1 + 0.5 * (1.7 - 1) = 1.35$$

$$W_{C3} = W_{C3} + \eta * (I_{13} - W_{C3})) = 1 + 0.5 * (1.8 - 1) = 1.4$$

Thus the updated weight for W_C is:

$$W_C(1) = [1.05, 1.35, 1.4]$$

We can similarly compute $W_A(1)$:

$$W_A(1) = [0.65, 1.2, 1.05]$$

Since B was not updated, $W_B(1) = W_B(0)$ (no change).

The new weight matrix $W(1)$ is then:

$$W(1) = \begin{pmatrix} \mathbf{w_A :} & \mathbf{0.65} & \mathbf{1.2} & \mathbf{1.05} \\ \mathbf{w_B :} & \mathbf{0.1} & \mathbf{0.1} & \mathbf{0.9} \\ \mathbf{w_C :} & \mathbf{1.05} & \mathbf{1.35} & \mathbf{1.4} \end{pmatrix}$$

For the next input $I_2(0, 0, 0)$:

We compute the distances of I_2 from A, B and C using the updated weight matrix from the last instance, ie. $W(1)$.

$$D^2(A, I_2) = 3, D^2(B, I_2) = 0.8, D^2(C, I_2) = 4.9$$

Since B is the closest, we pick B as the winning neuron and update weights of B and A.

$$W(2) = \begin{pmatrix} \mathbf{w_A :} & \mathbf{0.325} & \mathbf{0.6} & \mathbf{0.525} \\ \mathbf{w_B :} & \mathbf{0.05} & \mathbf{0.05} & \mathbf{0.45} \\ \mathbf{w_C :} & \mathbf{1.05} & \mathbf{1.35} & \mathbf{1.4} \end{pmatrix}$$

For the next input $I_3(0, 0.5, 1.5)$:

We compute the distances using $W(2)$, and get $D^2(A, I_2) = 1.1$, $D^2(B, I_2) = 1.3$, $D^2(C, I_2) = 1.7$

Hence we pick A as the winner, and update weights of A along with its neighbors with distance 1, which includes both B and C. (Thus in this case the weights of all 3 neurons get updated).

$$W(3) = \begin{pmatrix} \mathbf{w_A :} & \mathbf{0.16} & \mathbf{0.55} & \mathbf{1.01} \\ \mathbf{w_B :} & \mathbf{0.025} & \mathbf{0.275} & \mathbf{0.975} \\ \mathbf{w_C :} & \mathbf{0.525} & \mathbf{0.925} & \mathbf{1.45} \end{pmatrix}$$

Students can perform similar computations for I_4 , I_5 , and I_6 , to get the final weight matrix after the first epoch $W(6)$:

$$W(6) = \begin{pmatrix} w_A: & 0.77 & 0.69 & 0.75 \\ w_B: & 0.51 & 0.32 & 0.49 \\ w_C: & 0.76 & 0.86 & 0.99 \end{pmatrix}$$

Solution 3b):

To determine the clusters we compute the euclidean distances of the 6 inputs from the 3 neurons using the weight matrix $W(6)$, and assign the inputs to the cluster represented by their closest neuron.

The following table gives the values of distances and cluster assignments for all the inputs:

Input	Distance to A	Distance to B	Distance to C	Assigned Cluster
I_1	2.2	4.0	1.5	C
I_2	1.6	0.6	2.3	B
I_3	1.2	1.3	1.0	C
I_4	1.1	0.6	1.8	B
I_5	0.2	0.03	0.4	B
I_6	0.2	1.0	0.07	C

Solution 3c):

The learning rate for first epoch is $\eta(1) = 0.5$.

Since the rate decreases at a rate of 0.5, the learning rate for the second epoch is $\eta(2) = 0.5 \eta(1) = 0.5 * 0.5 = 0.25$

Similarly, the learning rate for the third epoch is $\eta(3) = 0.5 \eta(2) = 0.125$